

TCSS 558 — Applied Distributed Systems
Autumn 2015 — Project #4
Paxos-Based, Multi-threaded, Replicated Key-Value Store using RPC
Due Date: Midnight, Wednesday, December 9, 2015

Guidelines

Project should be electronically submitted to Canvas by midnight on the due date. A submission link is provided on the course Canvas page.

Assignment Overview

For this project, you will extend Project #3 by adding fault tolerance and achieving consensus of updates amongst replicated state machine KV-store servers using Paxos as described in the Lamport paper, "Paxos made simple" that you will deploy and test on the clustered system available through the University.

- 1) In Project #3 you replicated your Key-Value Store Server from Project #2 across 5 distinct servers and used 2PC to ensure consistency across replicas on KV-store operations (minimally PUT & DELETE). However, as we've discussed in class two-phase commit protocols are not fault tolerant. Your new goal for Project #4 is to integrate the capability to ensure continual operation of your KV-store despite replica failures. To achieve this goal you will implement Paxos to realize fault-tolerant consensus amongst your replicated servers. Functionally, you must implement and integrate the Paxos roles we described in class, and as described in the Lamport papers, including the Proposers, Acceptors, and Learners. The goal here is to focus on the Paxos implementation and algorithmic steps involved in realizing consensus in event ordering. Client threads may generate requests to any of the replicas at any time. To minimize the potential for live lock, you may choose to use leader election amongst the proposers, however, that is not a strict requirement of the project.
- 2) A second requirement for this project is that the acceptors must be configured to "fail" at random times. Each of the roles within Paxos may be implemented as threads or processes - that's up to you to determine how to implement (I'd use threads). Assuming you use threads for each role, at a minimum the acceptor threads should "fail" periodically, which could be done as simply as having a timeout that kills off the thread (or returns) after some random period of time. A new acceptor thread could then be restarted after another delay which should resume the functions of the previous acceptor thread, even though it clearly won't have the same state as the previously killed thread. Once this is completed, you may earn extra credit for the project if all roles are constructed to randomly fail and restart, but only the failure/restart of the acceptor is required. This should make it clear how Paxos overcomes replicated server failures.

As in project #1, you should use your client to pre-populate the Key-Value store with data and a set of keys. The composition of the data is up to you in terms of what you want to store there. Once the key-value store is populated, your client must do at least five of each operation: 5 PUTs, 5 GETs, 5 DELETEs.

Evaluation

Your Paxos-based, fault-tolerant replicated multi-threaded Key-Value Store servers will be evaluated on how well they inter-operate with each other using RPC while doing concurrent operations on the UWT-provided "cluster" as well as their conformance to the requirements above.

Executive Summary

Part of your completed assignment submission should be an executive summary containing an "assignment overview" (1 paragraph, up to about 250 words) explaining what you understand to be the purpose and scope of the assignment and a "technical impression" (1–2 paragraphs, about 200–500 words) describing your experiences while carrying out the assignment. The assignment overview shows how well you

understand the assignment; the technical impression section helps to determine what parts of the assignment need clarification, improvement, etc., for the future.

Grading

The grade for your executive summary is based on the effort you put into the assignment overview and technical impression. In general, if you put some effort into your writing, you will receive full credit for your executive summary.

Project Deliverables

The following items should be archived together, e.g., placed in a .zip file or tarball file (*.tgz or *.tar.gz), and electronically submitted via the link is provided on the course Canvas page.

- 1) All novel Java and/or C source code files implementing the two client and two server programs, i.e., plus any additional support code.
- 2) A simple README that includes
 - a) How to build your server and client codes (including any external libraries necessary)
 - b) How to run your server and client programs
- 3) Your executive summary